MA4199 Project – Bias Variance Tradeoff

Ng Wei Le

1 Approximation Theorem

Definition 1. The fill distance for a set of points $X = \{x_1, ..., x_N\} \subseteq \Omega$ for a bounded domain Ω is defined to be

$$h_{X,\Omega} \coloneqq \sup_{x \in \Omega} \min_{1 \leq j \leq N} \|x - x_j\|_2$$

.

The below theorem gives us some justification as to why the minimum norm interpolating function was chosen, though this only works under noiseless conditions:

Theorem 1. *Fix $h^* \in \mathcal{H}_\infty$. Let $(x_1, y_1), ..., (x_n, y_n)$ be i.i.d. random variables where x_i drawn randomly from a compact cube $\Omega \subseteq \mathbb{R}^d$, $y_i = h^*(x_i)\ \forall i$. There exists $A, B > 0$ such that for any interpolating $h \in \mathcal{H}_\infty$ with high probability*

$$\sup_{x \in \Omega} |h(x) - h^*(x)| < Ae^{-B(n/\log n)^{1/d}}(\|h^*\|_{\mathcal{H}_\infty} + \|h\|_{\mathcal{H}_\infty})$$

Theorem 11.22 in [6]:

Let Ω be a cube in $\mathbb{R}^d$. Suppose ... There exists a constant $c > 0$ such that the error between a function $f \in N(\Omega)$ and its interpolant $s_{f,X}$ can be bounded by:

$$\|f - s_{f,X}\|_{L_\infty(\Omega)} \leq \exp(-c/h_{X,\Omega})|f|_N(\Omega)$$

for all data sites X with sufficiently small $h_{X,\Omega}$.

With $h_{X,\Omega}$ as the fill on the order of $O(n/logn)^{-1/d}$ (using the theorem S1 in Belkin's paper which wasn't proved). We consider $f(x) := h(x) - h*(x)$. Since h is interpolating, we have $f(x_i) = 0$ for all x_i. We then let $s_{f,X}$ be the zero function, since it is an interpolant of f. Thus, we have: $s_{f,X}$ can be bounded by:

$$\|f\|_{L_\infty(\Omega)} = \sup_{x \in \Omega} |h(x) - h^*(x)| < \exp(-c(n/logn)^{1/d})|f|_N(\Omega)$$

$$\leq \exp(-c(n/logn)^{1/d})(\|h^*\|_{\mathcal{H}_\infty} + \|h\|_{\mathcal{H}_\infty})$$

Another form we can have is using proposition 14.1 in [6]:

Proposition 1. *Let $\Omega \subseteq \mathbb{R}^d$ be bounded and measurable. Suppose $X = \{x_1, ..., x_N\} \subseteq \Omega$ is quasi-uniform with respect to $c_{qu} > 0$. Then there exists constants $c_1, c_2 > 0$ depending only on space dimension d, on Ω and on c_{qu} such that:*

$$c_1 N^{-1/d} \leq h_{X,\Omega} \leq c_2 N^{-1/d}$$

.

With the definition of quasi-uniformness being:

Definition 2. For the separation distance of $X = \{x_1, ..., x_N\}$ being defined as $q_x \coloneqq \frac{1}{2} \min_{i \neq j} \|x_i - x_j\|_2$.

We can then use the above proposition with n replacing $n/\log n$.

In either case, by choosing a the smallest norm for h, we can see that it corresponds to the smallest upperbound for $|h(x) - h^*(x)|$.

2 Existing Bounds Provide No Guarantees for Interpolated Kernel Classifiers

Steps are:

- Find lower bound on function norm of t-overfitted classifiers in RKHS corresponding to Gaussian Kernels.
- Show loss for available bounds for kernel methods based on function norm (can perhaps use this to explain approximation theorem as well?)

Interpolation: 0 regression error. Overfitting: 0 classification error. Interpolation implies overfitting.

Definition 3. We say $h \in H$ t-overfits data, if it achieves zero classification loss (overfits) and $\forall_i y_i h(x_i) > t > 0$.

The below shows a theorem on how the function norm changes with respect to t-overfitting.

Theorem 2. *Let $(\mathbf{x}_i, y_i)$ be data sampled from P on $\Omega \times \{-1, 1\}$ for $i = 1, ..., n$. Assume that y is not a deterministic function of x on a subset of non-zero measure. Then, with high probability, any h that t-overfits the data, satisfies*

$$\|h\|_H > Ae^{Bn^{1/d}}$$

for some constants $A, B > 0$ depending on t.

We define the γ-shattering and fat-shattering dimension below:

Definition 4. Let F be a set of functions mapping from a domain X to $\mathbb{R}$. Suppose $S = \{x_1, x_2, ..., x_m\} \subseteq X$. Suppose also that γ is a positive real number. Then S is γ-shattered by F if there are real numbers $r_1, r_2, ..., r_m$, such that for each $b \in \{0, 1\}^m$ there is a function f_b in F with

$$f_b(x_i) \geq r_i + \gamma \text{ if } b_i = 1, \text{ and } f_b(x_i) \leq r_i - \gamma \text{ if } b_i = 0, \text{ for } 1 \leq i \leq m.$$

We say $r = (r_1, r_2, ..., r_m)$ witnesses the shattering. Suppose that F is a set of functions from a domain X to $\mathbb{R}$ and that $\gamma > 0$. Then F has γ-dimension d if d is the maximum cardinality of a subset S of X that is γ-shattered by F. If no such maximum exists, we say that F has infinite γ-dimension. The γ-dimension of F is denoted $fat_F(\gamma)$. This defines a function $fat_F : \mathbb{R} \to N \cup \{0, \infty\}$, which we call the fat-shattering dimension of F.

Proof. Let $B_R = \{f \in \mathcal{H}, \|f\|_{\mathcal{H}} < R\}$ be a ball of radius R in RKHS $\mathcal{H}$. Suppose the data is γ-overfitted, [3] gives us a high probability of a bound of

$$L(f) < O(\frac{\ln(n)^2}{\sqrt{n}}\sqrt{fat_{B_R}(\gamma/8)})$$

for $L(f)$ the expected classification error. Also, from [1] we have

$$fat_{B_R}(\gamma) < O((log(R/\gamma))^d)$$

. We then have B_R containing no function that γ overfits the data unless

$$(log(R/\gamma))^d > O(n) \implies R > c_1 \exp(c_2(\frac{n}{\ln n})^{1/d})$$

for some positive constants c_1, c_2. □

Classical bounds for kernel methods ([2]) are in the form:

$$|\frac{1}{n}\sum_i l(f(x_i), y_i) - L(f)| \leq C\frac{\|f\|_{\mathcal{H}}^a}{n^b}, \quad C, a, b \geq 0$$

The right side on this will tend to infinity for bigger $\|f\|_{\mathcal{H}}$, which is suggested by Theorem 2.

3 Random Fourier Features

For a feature map $\phi : \mathbb{R}^d \to \mathbb{R}^{d'}$ the kernel trick allows easy computation for positive definite kernel k where $k(x, y) =< \phi(x), \phi(y) >$. We want to find a randomized feature map $z : \mathbb{R}^d \to \mathbb{R}^{\bar{d}}$ such that

$$k(x, y) =< \phi(x), \phi(y) > \approx < z^{\mathrm{T}}(x), z(y) >$$

. As suggested by [4], for a shift-invariant kernel k: $k(x, y) = k(x - y)$), we consider the mapping $z(x) = cos(w^{\mathrm{T}}x + b)$, where w is drawn from the probability distribution p:

$$p(w) = \frac{1}{2\pi}\int k(h) \exp(-iw^{\mathrm{T}}h) \, \mathrm{d}h \tag{1}$$

when we compute the Fourier transform of the kernel k, and b is drawn from the uniform distribution on $[0, 2\pi]$.

We know that the fourier transform of $k(\cdot)$ is a probability distribution from Bochner's theorem:

Theorem 3. *(Bochner [5]).For a continuous kernel $k(x - y)$ it is a positive definite kernel if and only if $k(\cdot)$ is the fourier transform of a non-negative measure.*

We now have:

$$k(x - y) = \int_{\mathbb{R}^d} p(w)\exp(iw^{\mathrm{T}}(x - y)) \, \mathrm{d}w = \mathbb{E}_w[e^{iw^{\mathrm{T}}x}(e^{iw^{\mathrm{T}}y})^*]$$

. Therefore, we can use $e^{iw^{\mathrm{T}}x}(e^{iw^{\mathrm{T}}y})^*$ as an estimate (unbiased) of $k(x,y)$. Let $\phi_w(x) = e^{iw^{\mathrm{T}}x}$ We can also use $z_w(x) = \sqrt{2}cos(w^{\mathrm{T}}x + b)$ instead of $\phi_w(x)$, as suggested by [4].

Proposition 2. *For $z_w(x) = \sqrt{2}cos(w^T x + b)$, where w is drawn from probability distribution p in (1) and b drawn from a uniform random variable on $[0, 2\pi]$.*

$$E(z_w(x))z_w(y) = k(x,y)$$

Proof.

$$\begin{aligned} z_w(x) &= 2\ \frac{\sqrt{2}}{2}cos(w^{\mathrm{T}}x + b) \\ &= \frac{1}{\sqrt{2}}\ (e^{i(w^{\mathrm{T}}x+b)} + e^{-i(w^{\mathrm{T}}x+b)}) \\ &= \frac{1}{\sqrt{2}}\ (\phi_w(x)e^{ib} + \phi_w(x)^* e^{-ib}) \end{aligned}$$

Where $\phi_w(x) = e^{iw^{\mathrm{T}}x}$.

$$\begin{aligned} z_w(x)z_y(y) &= \frac{1}{2}[\phi_w(x)\phi_w(y)e^{i2b} + \phi_w(x)^*\phi_w(y)^* e^{-i2b} + \phi_w(x)\phi_w(y)^* + \phi_w(x)^*\phi_w(y)] \\ \mathbb{E}[z_w(x)z_y(y)] &= \frac{1}{2}\mathbb{E}[\phi_w(x)\phi_w(y)e^{i2b} + \phi_w(x)^*\phi_w(y)^* e^{-i2b}] + \frac{1}{2}\mathbb{E}[\phi_w(x)\phi_w(y)^*] + \frac{1}{2}\mathbb{E}[\phi_w(x)^*\phi_w(y)] \end{aligned}$$

As mentioned earlier in Theorem 3, $\mathbb{E}_w[\phi_w(x)\phi_w(y)^*] = k(x-y)$. Also $\phi_w(x)\phi_w(y)^* = (\phi_w(x)^*\phi_w(y))^*$.

$$\begin{aligned} \mathbb{E}[z_w(x)z_y(y)] &= \frac{1}{2}\mathbb{E}[\phi_w(x)\phi_w(y)e^{i2b} + \phi_w(x)^*\phi_w(y)^* e^{-i2b}] + \frac{1}{2}k(x-y) + \frac{1}{2}[k(x-y)]^* \\ &= \frac{1}{2}\mathbb{E}[\phi_w(x)\phi_w(y)e^{i2b} + \phi_w(x)^*\phi_w(y)^* e^{-i2b}] + k(x-y) \end{aligned}$$

For real kernel , $k(x-y) = (k(x-y))^*$.

$$\begin{aligned} \mathbb{E}_{w,b}[\phi_w(x)\phi_w(y)e^{i2b}] &= \frac{1}{2\pi}\int_{\mathbb{R}^d}\int_0^{2\pi} p(w)\phi_w(x)\phi_w(y)e^{i2b}\mathrm{d}b\ \mathrm{d}w \\ &= \frac{1}{2\pi}\int_{\mathbb{R}^d} p(w)\phi_w(x)\phi_w(y)\int_0^{2\pi} e^{i2b}\mathrm{d}b\ \mathrm{d}w \\ &= 0 \end{aligned}$$

Since $\int_0^{2\pi} e^{i2b}\mathrm{d}b = 0$. Similarly, $\mathbb{E}_{w,b}[\phi_w(x)^*\phi_w(y)^* e^{-i2b}] = 0$.

$$\therefore \mathbb{E}[z_w(x)z_y(y)] = k(x-y).$$

□

As suggested by [4], the variance of the estimate is decreased by using z, a D dimensional vector by

concatenating D of z_w and normalizing by a constant $\sqrt{D}$. We let:

$$z(x) = \sqrt{\frac{2}{D}}[cos(w_1^{\mathrm{T}}x + b_1)...cos(w_D^{\mathrm{T}}x + b_D)]$$

with randomly drawn w_i and b_i as described previously.

Theorem 4. *For N the number of random features, and $x_1, x_2, ..., x_n$ the data points, when $N > n$ and as N increases, the norm of the minimizer tends to the norm of the minimum norm RKHS interpolant.*

Proof. Let $f(x)$ be the minimum norm RKHS interpolant function for the datapoints.

$$f(x) = \sum_i \alpha_i k(x_i, x) \approx \sum_i \alpha_i z(x_i)^{\mathrm{T}} z(x) = \beta^{\mathrm{T}} z(x) = \hat{f}(x)$$

(the first equality holds due to Representer Theorem) Where $\beta = \sum_i \alpha_i z(x_i)$. The norm of the function from the random fourier features approximation is:

$$\|\beta\| = \beta^{\mathrm{T}}\bar{\beta} = (\sum_i \alpha_i z^{\mathrm{T}}(x_i))(\sum_i \bar{\alpha}_i \bar{z}(x_i)) = \sum_i \sum_j \alpha_i \bar{\alpha}_j z^{\mathrm{T}}(x_i)\bar{z}(x_j) \approx \sum_i \sum_j \alpha_i \bar{\alpha_j} k(x_i, x_j) = \|f\|$$

□

References

[1] Mikhail Belkin. Approximation beats concentration? an approximation view on inference with smooth radial kernels. *arxiv:1801.03437*, 2018.

[2] Mikhail Belkin, Siyuan Ma, and Soumik Mandal. To understand deep learning we need to understand kernel learning. *arxiv:1802.01396*, 2018.

[3] Balázs Kégl, Tamás Linder, and Gábor Lugosi. Data-dependent margin-based generalization bounds for classification. In *Lecture Notes in Computer Science*, pages 368–384. Springer Berlin Heidelberg, 2001.

[4] Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In J. Platt, D. Koller, Y. Singer, and S. Roweis, editors, *Advances in Neural Information Processing Systems*, volume 20, pages 1177–1184. Curran Associates, Inc., 2008.

[5] Walter Rudin. *Fourier Analysis on Groups*. John Wiley & Sons, Inc., jan 1990.

[6] Holger Wendland. *Scattered Data Approximation*. Cambridge Monographs on APplied and Computational Mathematics. Cambridge University Press, 2004.

Appendix